

THE NEGOTIATOR

15-Hour Roadmap — Residential Moving, 4-Person Team

Chosen vertical: residential moving. It's the brief's own worked example — benchmark data, red-flag rules, and the persona (Daniel, Rock Hill → Charlotte) are already provided, so zero research time is spent before building. The call list (Google Places) and the pain story (5.6x price spread, 40% overrun rate) are both ready to use on hour 0.

KEY DECISION	CHOICE
Vertical	Residential moving (B2C) — Rock Hill → Charlotte style route works out of the box.
Time budget	15 working hours, compressed into 6 phases below — tighter than a typical 20hr plan, so scope is trimmed hard.
Counterparty setup	Hybrid: real calls to 6-8 sourced movers via Twilio, plus 2 role-played calls (stonewaller, hard-sell upseller) as a guaranteed fallback for “3 distinct styles.”
Cut from scope	Document-intake path is built minimally (one document type, simple parsing) rather than fully polished; Lovable dashboard gets only the 3 must-have screens.

Team Roles

ROLE	OWNS
P1 — Voice & Conversation	Estimator voice interview + Caller agent conversation design (disclosure, friction handling, honesty line).
P2 — Telephony & Data	Twilio/SIP wiring, call list sourcing (Google Places), minimal document-intake parsing.
P3 — Backend & Negotiation Logic	Job-spec schema, config file, quote extraction, negotiation + red-flag logic, ranking engine.
P4 — Frontend & Demo	Lovable dashboard (3 screens only), demo script, backup recording.

Hour-by-Hour Roadmap (15 Hours)

Compressed from a 20-hour plan by cutting the document-intake path to one simple format, trimming the dashboard to 3 screens, and merging integration + hardening into one pass.

Phase 0 — Kickoff & Scoping		Hr 0-0.75
OWNER	TASK	
All	Confirm vertical (moving, already decided), freeze the job-spec schema, split into roles.	
P2	Register Twilio trial + phone number; pull 6-8 real mover numbers via Google Places for a test route.	
P3	Stub the config file: price benchmarks, 30%-below-market red-flag rule, negotiation levers.	
P4	Skeleton Lovable project: 3 empty screens (Confirm Spec / Live Calls / Report).	

Phase 1 — Foundation Build		Hr 0.75-3
OWNER	TASK	
P1	Build the Estimator voice agent (ElevenLabs): asks rooms, large items, stairs, access, date — outputs the frozen schema.	
P2	Minimal document intake: one format only (e.g. a photo of an inventory list) parsed via OpenAI vision into the same schema.	
P3	Stand up the backend API: job-spec CRUD + call/quote logging endpoints.	
P4	Build the “Confirm your spec” screen against the real backend.	

Phase 2 — The Caller		Hr 3-6
OWNER	TASK	
P1	Build the Caller agent's script: AI disclosure line, consistent job description, itemized-quote ask, graceful handling of “robot?” and “no phone quotes.”	
P2	Wire ElevenLabs → Twilio; get one full test call connecting end-to-end.	
P3	Build the quote-extraction tool: transcript → structured itemized quote.	
P4	Build the “Live Calls” screen: status per company + playback once available.	
Checkpoint	One full call end-to-end before Hr 6 — non-negotiable gate.	

Phase 3 — Parallel Quote Gathering		Hr 6-8.5
OWNER	TASK	
P2	Fire the real batch to the 6-8 sourced movers.	
P1 + P4 (role-play)	Run the 2 backup calls: one stonewaller, one hard-sell upseller — guarantees 3 negotiation styles regardless of real pickup rate.	

OWNER	TASK
P3	Log every returned call as a structured, itemized quote tagged by negotiation style.
Checkpoint	At least 3 structured quotes by Hr 8.5 — prioritize this over any polish.

Phase 4 — The Closer: Negotiate & Report

Hr 8.5-11.5

OWNER	TASK
P3	Build negotiation logic: re-call cites a competing quote, pushes on fees, flags <70%-of-median quotes as warnings.
P1	Run the actual re-call that moves price/terms — record it, this is the demo centerpiece.
P3	Build ranking + plain-language recommendation with transcript citations.
P4	Build the final “Report” screen: ranked table, recommendation banner, itemized fees, red-flag callouts.

Phase 5 — Integration & Hardening

Hr 11.5-13

OWNER	TASK
All	One full run-through: intake → confirm → calls → negotiation → report. Fix seams only, no new features.
P1	Verify the 4 conversation requirements (disclosure, friction, honesty, structured endings) against real recordings.
P3	Confirm the config file is genuinely swappable (rename one field as a test).

Phase 6 — Demo Prep

Hr 13-15

OWNER	TASK
P4	Write the 3-minute demo script: problem → spec confirm → call clips → negotiation moment → ranked report.
All	Record a full backup video in case live calls or Wi-Fi fail during judging.
All	Rehearse twice, timed — cut anything outside intake → calls → negotiation → report.

Final Check Before Demo

- Loop closed end-to-end: intake → calls → negotiation → ranked recommendation with transcript evidence.
- One structured job spec, built by voice interview **and** the one document type, confirmed by the user.
- Live calls against at least 3 distinct negotiation styles (real + role-played), every quote itemized.
- At least one negotiation shows price or terms measurably change mid-call because of gathered leverage.
- AI disclosure and honesty constraints hold throughout every recorded call.
- Every call ends in a structured outcome, never a vague number.
- Final report ranks all quotes, cites recordings/transcripts, explains the recommendation in plain language.

What Got Cut to Fit 15 Hours

- Document intake supports one format only (not multiple types) — still satisfies the “at least one document type” requirement.
- Dashboard is 3 screens, no extra polish pass — report can fall back to a clean printout with audio clips if Lovable runs behind.
- No separate hardening phase for edge cases beyond the 4 conversation requirements — fix only what's demoed.
- Skip building counter-agents for agent-to-agent simulation; role-play covers the negotiation-style requirement faster.